



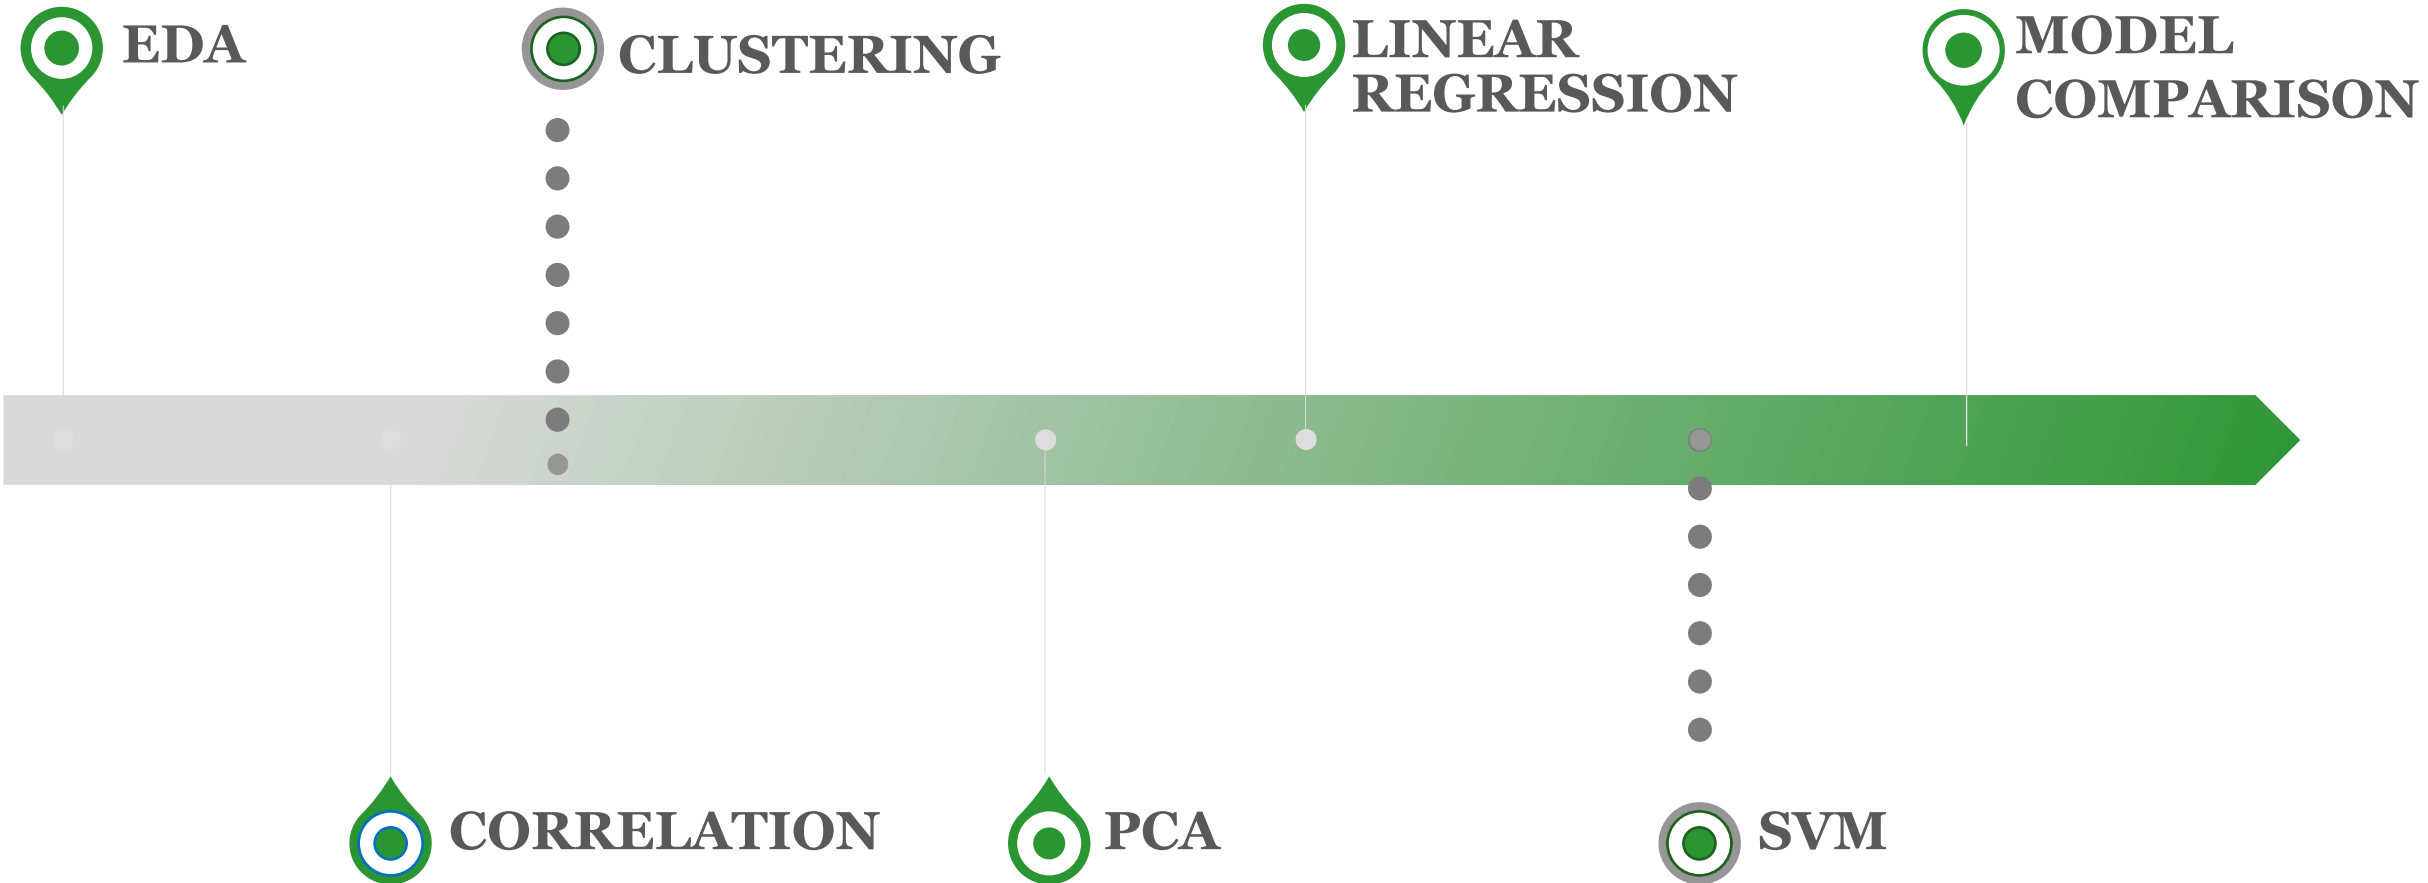
# STUDENTS PERFORMANCE PREDICTION

---

Nava Carlo  
Passaro Jacopo



# Agenda



# Dataset


	school	sex	age	address	famsize	Pstatus	Medu	Fedu	Mjob	Fjob	reason	guardian	traveltime	studytime	failures	schoolsup	famsup	paid
1	GP	F	18	U	GT3	A	4	4	at_home	teacher	course	mother	2	2	0	yes	no	no
2	GP	F	17	U	GT3	T	1	1	at_home	other	course	father	1	2	0	no	yes	no
3	GP	F	15	U	LE3	T	1	1	at_home	other	other	mother	1	2	3	yes	no	yes
4	GP	F	15	U	GT3	T	4	2	health	services	home	mother	1	3	0	no	yes	yes
5	GP	F	16	U	GT3	T	3	3	other	other	home	father	1	2	0	no	yes	yes
6	GP	M	16	U	LE3	T	4	3	services	other	reputation	mother	1	2	0	no	yes	yes
7	GP	M	16	U	LE3	T	2	2	other	other	home	mother	1	2	0	no	no	no
8	GP	F	17	U	GT3	A	4	4	other	teacher	home	mother	2	2	0	yes	yes	no
9	GP	M	15	U	LE3	A	3	2	services	other	home	mother	1	2	0	no	yes	yes
10	GP	M	15	U	GT3	T	3	4	other	other	home	mother	1	2	0	no	yes	yes
11	GP	F	15	U	GT3	T	4	4	teacher	health	reputation	mother	1	2	0	no	yes	yes

activities	nursery	higher	internet	romantic	famrel	freetime	goout	Dalc	Walc	health	absences	G1	G2	G3
no	yes	yes	no	no	4	3	4	1	1	3	6	5	6	6
no	no	yes	yes	no	5	3	3	1	1	3	4	5	5	6
no	yes	yes	yes	no	4	3	2	2	3	3	10	7	8	10
yes	yes	yes	yes	yes	3	2	2	1	1	5	2	15	14	15
no	yes	yes	no	no	4	3	2	1	2	5	4	6	10	10
yes	yes	yes	yes	no	5	4	2	1	2	5	10	15	15	15
no	yes	yes	yes	no	4	4	4	1	1	3	0	12	12	11
no	yes	yes	no	no	4	1	4	1	1	1	6	6	5	6
no	yes	yes	yes	no	4	2	2	1	1	1	0	16	18	19

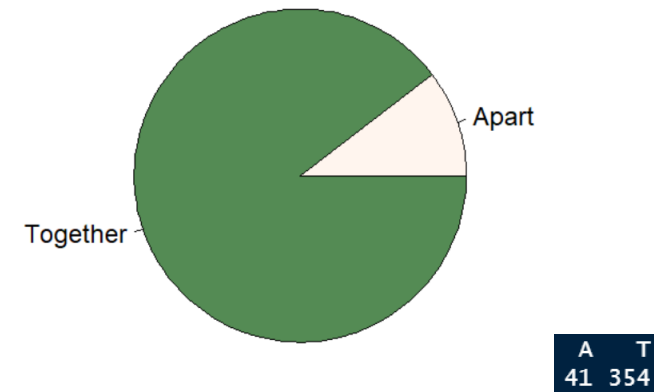
Variables: 33  
Observations: 395  
Target: G3

# Exploratory Data Analysis


## Gender Distribution




## Parent Status



## Father Occupation




## Mother Occupation




# Exploratory Data Analysis


## Father vs Mother Occupation




## Family Relations



## Mother Job Distrib. Per Education




## Father Job Distrib. Per Education

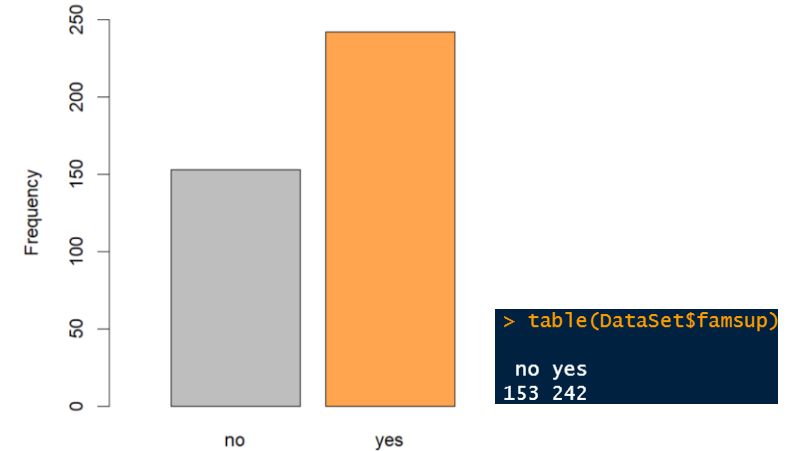


# Exploratory Data Analysis

## Father vs Mother Education



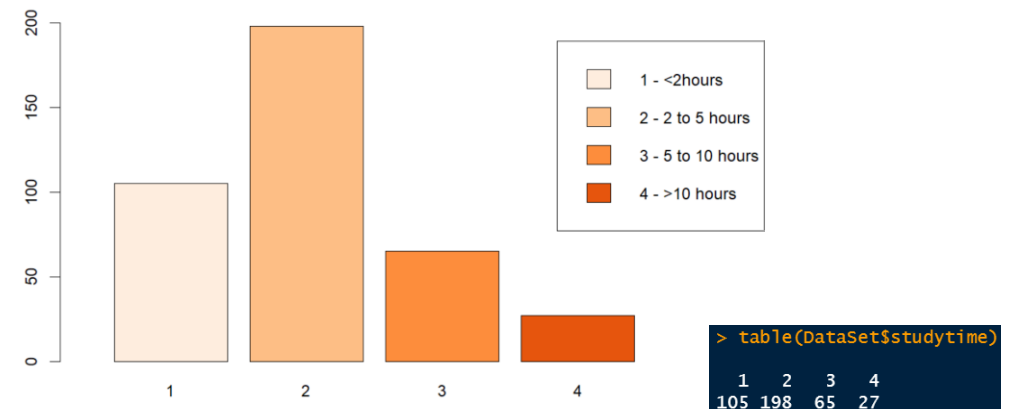
## Family Support



## Failures – Family Support Matrix


```
> table(DataSet[,c(17,15)])
      failures
famsup 0  1  2  3
no    115 25  5  8
yes   197 25 12  8
```

## Weekly Studytime




# Exploratory Data Analysis

## Count of Past Failures




## Grade 1 Distribution



```
summary(DataSet$G1)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
3.00	8.00	11.00	10.91	13.00	19.00


## Grade 2 Distribution



```
summary(DataSet$G2)
```

Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.00	9.00	11.00	10.71	13.00	19.00

## Grade 3 Distribution




```
summary(DataSet$G3)
```


Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
0.00	8.00	11.00	10.42	14.00	20.00

# Correlation

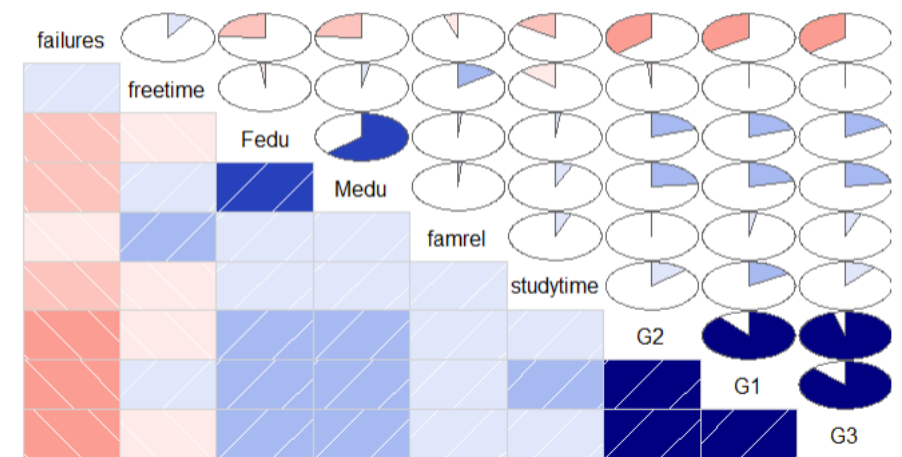
## Correlation Matrix - Pearson



## Correlation Matrix - Kendall




## Correlation Matrix - Spearman





# Clustering

Cluster Dendrogram



Variable Selected for the Cluster analysis

```
subset.clus <- subset(DataSet, select = c(Medu, Fedu, famrel, G1, G2, G3))
subset.sc <- scale(subset.clus)
subset.dist <- dist(subset.sc)
```

```
> subset.hc
```

```
Call:
hclust(d = subset.dist, method = "complete")
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 395
```

Distribution of G1, G2, G3  
in the Two Clusters


```
> table(subset.clus$G1, subset.hc.2)
subset.hc.2
  1  2
3  1  0
4  1  0
5  7  0
6 24  0
7 37  0
8 41  0
9 28  3
10 39 12
11 16 23
12 11 24
13  3 30
14  3 27
15  0 24
16  0 22
17  0  8
18  0  8
19  0  3
```

```
> table(subset.clus$G2, subset.hc.2)
subset.hc.2
  1  2
0 13  0
4  1  0
5 15  0
6 14  0
7 21  0
8 32  0
9 48  2
10 38  8
11 12 23
12 12 29
13  5 32
14  0 23
15  0 34
16  0 13
17  0  5
18  0 12
19  0  3
```


```
> table(subset.clus$G3, subset.hc.2)
subset.hc.2
  1  2
0 38  0
4  1  0
5  7  0
6 15  0
7  9  0
8 32  0
9 26  2
10 45 11
11 21 26
12 11 20
13  4 27
14  2 25
15  0 33
16  0 16
17  0  6
18  0 12
19  0  5
20  0  1
```

# Clustering

Distribution of G1, G2, G3 in the 1<sup>st</sup> Cluster



Distribution of G1, G2, G3 in the 2<sup>nd</sup> Cluster



# Principal Component Analysis

## Convert some Categorical variables into Numerical

```
sex_bin <- ifelse(DataSet$sex == "M", 1, 0)
Pstatus_bin <- ifelse(DataSet$Pstatus == "T", 1, 0)
famsup_bin <- ifelse(DataSet$famsup == "yes", 1, 0)
schoolsup_bin <- ifelse(DataSet$schoolsup == "yes", 1, 0)
romantic_bin <- ifelse(DataSet$romantic == "yes", 1, 0)
```

```
## Replacing zeros in G2 and G3 with NA
Subset$G2 <- ifelse(Subset$G2==0, NA, Subset$G2)
Subset$G3 <- ifelse(Subset$G3==0, NA, Subset$G3)
## Omitting NAs
Subset <- na.omit(Subset)
```

## Apply the PCA algorithm to the numerical Subset

```
> pca = prcomp(Subset, scale = TRUE)
> summary(pca)
```

Importance of components:


	PC1	PC2	PC3	PC4	PC5	PC6	PC7	PC8
Standard deviation	1.7988	1.4793	1.30091	1.15165	1.05276	1.02990	1.0099	0.9229
Proportion of Variance	0.1903	0.1287	0.09955	0.07802	0.06519	0.06239	0.0600	0.0501
Cumulative Proportion	0.1903	0.3191	0.41862	0.49663	0.56183	0.62422	0.6842	0.7343

	PC9	PC10	PC11	PC12	PC13	PC14	PC15
Standard deviation	0.91189	0.89207	0.81883	0.81330	0.77870	0.60715	0.57297
Proportion of Variance	0.04891	0.04681	0.03944	0.03891	0.03567	0.02168	0.01931
Cumulative Proportion	0.78323	0.83004	0.86948	0.90839	0.94406	0.96575	0.98506

	PC16	PC17
Standard deviation	0.41514	0.2857
Proportion of Variance	0.01014	0.0048
Cumulative Proportion	0.99520	1.0000



# Linear Regression

## Linear Regression on all the features

```
lm0 <- lm(DataSet.G3 ~ ., data=Subset)
summary(lm0)
```

```
Call:
lm(formula = DataSet.G3 ~ ., data = Subset)
```

```
Residuals:
    Min     1Q   Median     3Q      Max
-9.1042 -0.5100  0.3037  0.9716  4.2735
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  -3.12840    0.81082  -3.858 0.000134 ***
sex_bin       0.08746    0.21712   0.403 0.687308
Pstatus_bin  -0.26402    0.31688  -0.833 0.405271
DataSet.Medu  0.11919    0.11583   1.029 0.304137
DataSet.Fedu  -0.16640    0.11505  -1.446 0.148922
famsup_bin    0.22552    0.20575   1.096 0.273744
DataSet.famrel 0.33154    0.10892   3.044 0.002498 **
DataSet.studytime -0.16523    0.12555  -1.316 0.188979
DataSet.failures -0.26186    0.14360  -1.824 0.069006 .
schoolsup_bin  0.54869    0.29901   1.835 0.067289 .
romantic_bin  -0.32589    0.20887  -1.560 0.119527
DataSet.freetime 0.05040    0.10141   0.497 0.619490
DataSet.Dalc  -0.15501    0.14302  -1.084 0.279134
DataSet.Walc   0.17907    0.10083   1.776 0.076535 .
DataSet.absences 0.03680    0.01243   2.960 0.003265 **
DataSet.G1     0.18883    0.05790   3.262 0.001209 **
DataSet.G2     0.95510    0.04985  19.161 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.88 on 378 degrees of freedom
Multiple R-squared:  0.8385,    Adjusted R-squared:  0.8317
F-statistic: 122.7 on 16 and 378 DF, p-value: < 2.2e-16
```

## Linear Regression on grades G1 and G2

```
lm1 <- lm(DataSet.G3 ~ DataSet.G1 + DataSet.G2, data=Subset)
summary(lm1)
```

```
Call:
lm(formula = DataSet.G3 ~ DataSet.G1 + DataSet.G2, data = Subset)
```

```
Residuals:
    Min     1Q   Median     3Q      Max
-9.5713 -0.3888  0.2885  0.9725  3.7089
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)  -1.83001    0.33531  -5.458 8.57e-08 ***
DataSet.G1     0.15327    0.05618   2.728 0.00665 **
DataSet.G2     0.98687    0.04957  19.909 < 2e-16 ***
---
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 1.937 on 392 degrees of freedom
Multiple R-squared:  0.8222,    Adjusted R-squared:  0.8213
F-statistic: 906.1 on 2 and 392 DF, p-value: < 2.2e-16
```

# Linear Regression

## Linear Regression on family variables

```
lm2 <- lm(DataSet.G3 ~ Pstatus_bin + DataSet.Fedu + DataSet.Medu + famsup_bin + DataSet.famrel, data=Subset)
summary(lm2)
```

```
Call:
lm(formula = DataSet.G3 ~ Pstatus_bin + DataSet.Fedu + DataSet.Medu +
    famsup_bin + DataSet.famrel, data = Subset)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-12.2716  -2.0376   0.6069   2.9067   9.6411
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)    7.4851     1.3934   5.372 1.34e-07 ***
Pstatus_bin   -0.4354     0.7451  -0.584 0.55932
DataSet.Fedu    0.1551     0.2662   0.583 0.56056
DataSet.Medu    0.8617     0.2656   3.245 0.00128 **
famsup_bin    -0.7723     0.4728  -1.633 0.10323
DataSet.famrel  0.2620     0.2516   1.041 0.29842
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 4.475 on 389 degrees of freedom
Multiple R-squared:  0.05786,    Adjusted R-squared:  0.04575
F-statistic: 4.778 on 5 and 389 DF,  p-value: 0.000299
```

## Linear Regression on studytime and failures

```
lm3 <- lm(DataSet.G3 ~ DataSet.studytime + DataSet.failures, data=Subset)
summary(lm3)
```

```
Call:
lm(formula = DataSet.G3 ~ DataSet.studytime + DataSet.failures,
    data = Subset)
```

```
Residuals:
    Min       1Q   Median       3Q      Max
-11.5342  -1.9556   0.0613   3.0359   9.2429
```

```
Coefficients:
            Estimate Std. Error t value Pr(>|t|)
(Intercept)    10.7402     0.5970  17.991 < 2e-16 ***
DataSet.studytime  0.1985     0.2610   0.761  0.447
DataSet.failures  -2.1815     0.2945  -7.407 7.97e-13 ***
---
```

```
Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

```
Residual standard error: 4.281 on 392 degrees of freedom
Multiple R-squared:  0.1312,    Adjusted R-squared:  0.1267
F-statistic: 29.59 on 2 and 392 DF,  p-value: 1.072e-12
```

By comparing the Adjusted R-squared of the models we can assess that the best one is the linear regression on all the features.

# Support Vector Machine

## SVM on the whole Train Data

```
> subset.lsvm <- svm(DataSet.G3 ~., data = train.svm, type = "C-classification", kernel = "linear")
```

```
> pred <- predict(subset.lsvm, test.svm)
> table(pred, DataSet.G3)
      DataSet.G3
pred 0  5  6  7  8  9 10 11 12 13 14 15 16 17 18
 0    4  0  0  1  0  1  0  0  0  0  0  0  0  0  0
 4    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 5    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 6    0  0  2  0  2  0  0  0  0  0  0  0  0  0  0
 7    0  1  0  0  0  0  2  0  0  0  0  0  0  0  0
 8    1  1  2  1  1  0  0  1  0  0  0  0  0  0  0
 9    2  0  0  1  2  2  2  0  0  0  0  0  0  0  0
10    1  0  0  0  1  2  3  1  0  0  0  0  0  0  0
11    1  0  0  0  0  1  0  1  3  1  1  0  0  0  0
12    0  0  0  0  0  0  1  2  0  0  0  0  0  0  0
13    0  0  0  0  0  0  0  2  3  0  6  0  0  0  0
14    0  0  0  0  0  0  0  0  2  0  2  1  0  0  0
15    0  0  0  0  0  0  0  0  0  0  3  3  4  0  1
16    0  0  0  0  0  0  0  0  0  0  0  1  1  1  1
17    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
18    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
19    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
20    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
```

Make the prediction  
of G3 on Test Data

## Set Train and Test Partition

```
sample.svm <- sample(nrow(Substet), nrow(Substet)*0.8)
train.svm <- Substet[sample.svm,]
train.svm <- data.frame(train.svm)

test.svm <- Substet[-sample.svm,]
test.svm <- data.frame(test.svm)
```

## SVM on Train Data – Related only to family features

```
> fam.lsvm <- svm(DataSet.G3 ~ Pstatus_bin + DataSet.Medu + DataSet.Fedu + DataSet.famrel
+ famsup_bin + schoolsup_bin, data = train.svm, type = "C-classification", kernel = "linear")
```

Make the prediction  
of G3 on Test Data

```
> table(fam.pred, DataSet.G3)
      DataSet.G3
fam.pred 0  5  6  7  8  9 10 11 12 13 14 15 16 17 18
 0    1  0  0  0  0  0  0  1  0  0  0  1  0  0  0
 4    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 5    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 6    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 7    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 8    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
 9    0  0  1  0  0  0  0  0  0  0  0  0  0  0  0
10    8  2  2  2  5  5  6  5  7  0  9  4  3  1  0
11    0  0  1  0  1  0  0  1  1  1  2  0  1  0  0
12    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
13    0  0  0  1  0  0  2  0  0  0  0  0  1  0  0
14    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
15    0  0  0  0  0  1  0  0  0  0  1  0  0  0  2
16    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
17    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
18    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
19    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
20    0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
```



# Support Vector Machine

### SVM on Train Data – Related only to study and alcohol features

```
> stu.lsvm <- svm(DataSet.G3 ~ DataSet.studytime + DataSet.failures + DataSet.freetime +
  DataSet.Dalc + DataSet.Walc + DataSet.absences, data = train.svm, type = "C-classification", kernel = "linear")
```

```
> stu.pred <- predict(stu.lsvm, test.svm)
> table(stu.pred, DataSet.G3)
      DataSet.G3
stu.pred  0  5  6  7  8  9 10 11 12 13 14 15 16 17 18
0      7  0  0  0  0  1  2  0  0  0  1  1  0  0  0
4      0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
5      0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
6      0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
7      0  0  0  0  0  0  1  0  0  0  0  0  0  0  0
8      0  0  0  1  1  1  1  2  0  0  0  0  0  0  0
9      0  0  0  1  0  0  0  0  0  0  0  0  0  0  0
10     2  1  2  0  2  1  2  3  5  1  5  3  2  0  2
11     0  0  2  1  2  2  0  1  2  0  4  1  1  0  0
12     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
13     0  1  0  0  0  1  1  0  0  0  0  0  0  0  0
14     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
15     0  0  0  0  1  0  1  1  1  0  2  0  2  1  0
16     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
17     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
18     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
19     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
20     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
```

## Make the prediction of G3 on Test Data

### SVM on Train Data – Related only to failures, G1, G2

```
> grade1svm <- svm(DataSet.G3 ~ DataSet.failures + DataSet.G1 + DataSet.G2, data = train.svm, type = "C-classification", kernel = "linear")
```


## Make the prediction of G3 on Test Data

```
> grade.pred <- predict(stu.lsvm, test.svm)
> table(grade.pred, DataSet.G3)
      DataSet.G3
grade.pred  0  5  6  7  8  9 10 11 12 13 14 15 16 17 18
0      7  0  0  0  0  0  1  2  0  0  0  1  1  0  0
4      0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
5      0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
6      0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
7      0  0  0  0  0  0  1  0  0  0  0  0  0  0  0
8      0  0  0  1  1  1  1  2  0  0  0  0  0  0  0
9      0  0  0  1  0  0  0  0  0  0  0  0  0  0  0
10     2  1  2  0  2  1  2  3  5  1  5  3  2  0  2
11     0  0  2  1  2  2  0  1  2  0  4  1  1  0  0
12     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
13     0  1  0  0  0  1  1  0  0  0  0  0  0  0  0
14     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
15     0  0  0  0  1  0  1  1  1  0  2  0  2  1  0
16     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
17     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
18     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
19     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
20     0  0  0  0  0  0  0  0  0  0  0  0  0  0  0
```


# Support Vector Machine

**Plot of Predicted Values against Real Values for all the previous models**


**Predicted G3 vs Real G3**




**Predicted G3 vs Real G3 : FAM SVM**



**Predicted G3 vs Real G3 : STU SVM**



**Predicted G3 vs Real G3 : GRADE SVM**





# Model Comparison

## Linear Regression on G1, G2 vs SVM on failures, G1, G2

```
pred.lm1 <- predict(lm1, newdata = test.svm)
pred.grade.lsvm <- predict(grade.lsvm, newdata = test.svm)
```

```
rmse_lm1 <- rmse(DataSet.G3, pred.lm1)
rmse_grade.lsvm <- rmse(DataSet.G3, as.numeric(pred.grade.lsvm))
```

```
> rmse_lm1
[1] 2.377473
> rmse_grade.lsvm
[1] 3.05436
```

→ Linear Regression wins

## Linear Regression on Family variables vs SVM on Family features

```
pred.lm2 <- predict(lm2, newdata = test.svm)
pred.fam.lsvm <- predict(fam.lsvm, newdata = test.svm)
```

```
rmse_lm2 <- rmse(DataSet.G3, pred.lm2)
rmse_fam.lsvm <- rmse(DataSet.G3, as.numeric(pred.fam.lsvm))
```

```
> rmse_lm2
[1] 4.620417
> rmse_fam.lsvm
[1] 5.112655
```

Linear Regression wins ←

## Linear Regression on the whole Dataset vs SVM on the whole Dataset

```
pred.lm0 <- predict(lm0, newdata = test.svm)
pred.subset.lsvm <- predict(subset.lsvm, newdata = test.svm)
```

```
rmse_lm0 <- rmse(DataSet.G3, pred.lm1)
rmse_subset.lsvm <- rmse(DataSet.G3, as.numeric(pred.grade.lsvm))
```

```
> rmse_lm0
[1] 2.377473
> rmse_subset.lsvm
[1] 3.05436
```

→ Linear Regression wins

Thanks for your  
Attention

Nava Carlo  
Passaro Jacopo

---